

Random Walk Metropolis-Hastings Algorithm

Ying-Ting Lin

April 1, 2016

Goal:

Build a Random Walk Metropolis-Hastings function to derive the posterior marginal densities of parameters in a Gamma distribution with Jeffery's prior.

Jeffery's prior for $\text{Gamma}(\alpha, \beta)$ distribution:

$$\pi_J(\theta) \propto \sqrt{\det(I(\theta))} = \sqrt{\text{trigamma}(\alpha) \frac{\alpha}{\beta^2} - [\frac{1}{\beta}]^2} = \frac{\sqrt{\alpha * \text{trigamma}(\alpha) - 1}}{\beta}$$

Random Walk Metropolis-Hastings:

- Step 1: generate some observed data
- Step 2: choose the starting values alpha & beta
- Step 3: use Gibbs Sampling to generate a new alpha & beta
- Step 4: compute acceptance probability
- Step 5: accept or not

```
GammaMHExample <- function(n.sim, n.burnin){  
  library(MASS)  
  alpha <- 3; beta <- 3 #pick the parameter values to generate a random "observed" data  
  N <- 30 #set observed sample size  
  x <- rgamma(N, shape = alpha, scale = 1/beta) #generate observed data of size N  
  
  #defines a function that returns unnormalized posterior density  
  post_den <- function(alpha, beta, x, N){  
    #Jeffery's prior  
    prior <- sqrt(alpha*trigamma(alpha)-1)/beta  
    #gamma distr  
    likelihood <- prod(x^(alpha-1))*prod(exp(-beta*x))*((beta^alpha)/gamma(alpha))^N  
    # unnormalized target distr  
    working_cond_den <- prior*likelihood  
  }  
  
  theta.mh <- matrix(NA, nrow = n.sim, ncol = 2) #empty matrix to store sampled parameters  
  rho <- 0.05  
  #generate inital parameters  
  theta.current <- mvrnorm(n=1, mu = c(3,3), Sigma = matrix(c(0.3,rho,rho,0.6), nrow=2, ncol=2))  
  theta.update <- function(index, theta.current, x, N, i, rho) {  
    if (index == 1){ #if it's sample for alpha, update the first element of the parameter vector  
      #sample a new parameter value from Gaussian  
      theta.star <- rnorm(n = 1, mean = theta.current[index], sd = sqrt(0.3-rho^2))  
      #resample is drawn value is not positive  
      while(theta.star <= 0){  
        theta.star <- rnorm(n = 1, mean = theta.current[index], sd = sqrt(0.6-rho^2))  
      }  
    }  
  }  
}
```

```

    theta.temp <- c(theta.star, theta.current[2])
  }
  else {
    theta.star <- rnorm(n = 1, mean = theta.current[index]+
      rho*(theta.current[1]-theta.mh[i, 1]), sd = sqrt(1-rho^2))
    #resample is drawn value is not positive
    while(theta.star <=0){
      theta.star <- rnorm(n = 1, mean = theta.current[index]+
        rho*(theta.current[1]-theta.mh[i, 1]), sd = sqrt(1-rho^2))
    }
    theta.temp <- c(theta.current[1], theta.star)#for beta, update the second element
  }
  #compute MH ratio
  r <- post_den(theta.temp[1],theta.temp[2],x,N)/
    post_den(theta.current[1],theta.current[2],x,N)
  r <- min(r, 1, na.rm = TRUE) # r is the acceptance probability, NA values are ignored
  if (runif(1) < r)
    #runif(1) generates a uniform random number bewteen 0 and 1
    #if runif(1) < r, accept the new value; else reject the new value and keep the old value
    theta.star
  else theta.current[index]
}
for (i in 1:n.sim) {
  #iteration for alpha
  theta.current[1] <- theta.mh[i, 1] <- theta.update(1, theta.current,x, N, i,rho)
  #iteration for beta
  theta.current[2] <- theta.mh[i, 2] <- theta.update(2, theta.current,x, N, i,rho)
}
theta.mh <- theta.mh[(n.burnin + 1):n.sim, ] #discard burn-in
}

mh.draws <- GammaMHExample(n.sim = 10000, n.burnin = 1000)

```

Summary Statistics

```

library(coda)
mh.draws <- mcmc(mh.draws)
summary(mh.draws)

```

```





##
## Iterations = 1:9000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 9000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##      Mean      SD Naive SE Time-series SE

```

```
## [1,] 2.246 0.5352 0.005641      0.03574
## [2,] 2.164 0.5723 0.006033      0.03596
##
## 2. Quantiles for each variable:
##
##      2.5%  25%  50%  75% 97.5%
## var1 1.374 1.857 2.195 2.586 3.454
## var2 1.200 1.742 2.099 2.513 3.426
```


Diagnostics

```
#Traceplot
plot(mh.draws)
```



The traceplots suggest a good mixing, since the sample draws stay within a reasonable range of the parameter space and do not stuck in any particular area.

```
#Autocorrelation function plots
autocorr.plot(mh.draws)
```



The autocorrelation plots show that the autocorrelation is decreasing but still stays positive for a long lag.

#Gelman and Rubin

```
mh.draws1 <- mcmc(GammaMHExample(n.sim = 10000, n.burnin = 1000))
mh.draws2 <- mcmc(GammaMHExample(n.sim = 10000, n.burnin = 1000))
mh.draws3 <- mcmc(GammaMHExample(n.sim = 10000, n.burnin = 1000))
mh.draws4 <- mcmc(GammaMHExample(n.sim = 10000, n.burnin = 1000))
mh.draws5 <- mcmc(GammaMHExample(n.sim = 10000, n.burnin = 1000))
mh.list <- mcmc.list(list(mh.draws1, mh.draws2, mh.draws3, mh.draws4, mh.draws5))
gelman.diag(mh.list)
```

```
## Potential scale reduction factors:
```

```
##
```

```
##      Point est. Upper C.I.
```

```
## [1,]      1.07      1.18
```

```
## [2,]      1.35      1.79
```

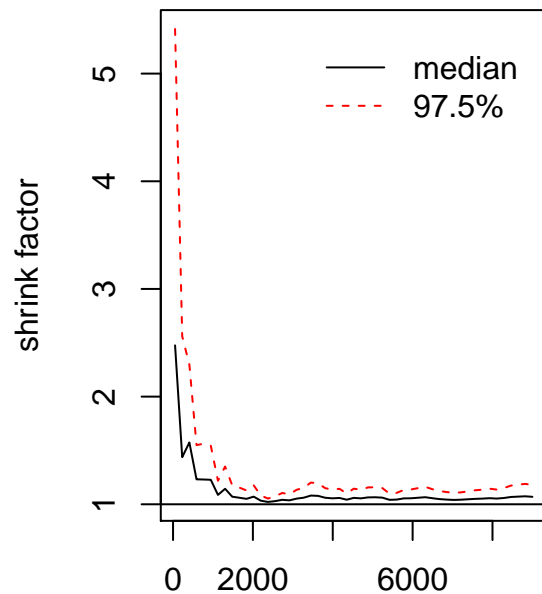
```
##
```

```
## Multivariate psrf
```


```
##
```

```
## 2.02
```

```
gelman.plot(mh.list)
```



last iteration in chain



last iteration in chain

The scale reduction factors reduces greatly after the first 1,000 iterations, and the median scale reduction factors are not far away from 1. The results suggest that between-chain variance is not much greater than within-chain variance and that the chains have converged to the stationary distribution with 10,000 sample draws.

```
#Geweke
geweke.diag(mh.draws)
```

```
##
## Fraction in 1st window = 0.1
## Fraction in 2nd window = 0.5
##
## var1 var2
## 1.305 1.275
```

The Z-scores are low, and hence we cannot reject the null hypothesis that fractions in the two windows are from the same distribution.

```
#Raftery and Lewis
raftery.diag(mh.draws,q=0.025,r=0.005,s=0.95)
```

```
##
## Quantile (q) = 0.025
## Accuracy (r) = +/- 0.005
## Probability (s) = 0.95
##
## Burn-in Total Lower bound Dependence
## (M) (N) (Nmin) factor (I)
## 23 24843 3746 6.63
## 29 31539 3746 8.42
```

Dependence factors exceed the benchmark 5, which may be caused from influential starting values, high correlations between parameters, or poor mixing.

```
#Heidelberg and Welch  
heidel.diag(mh.draws)
```

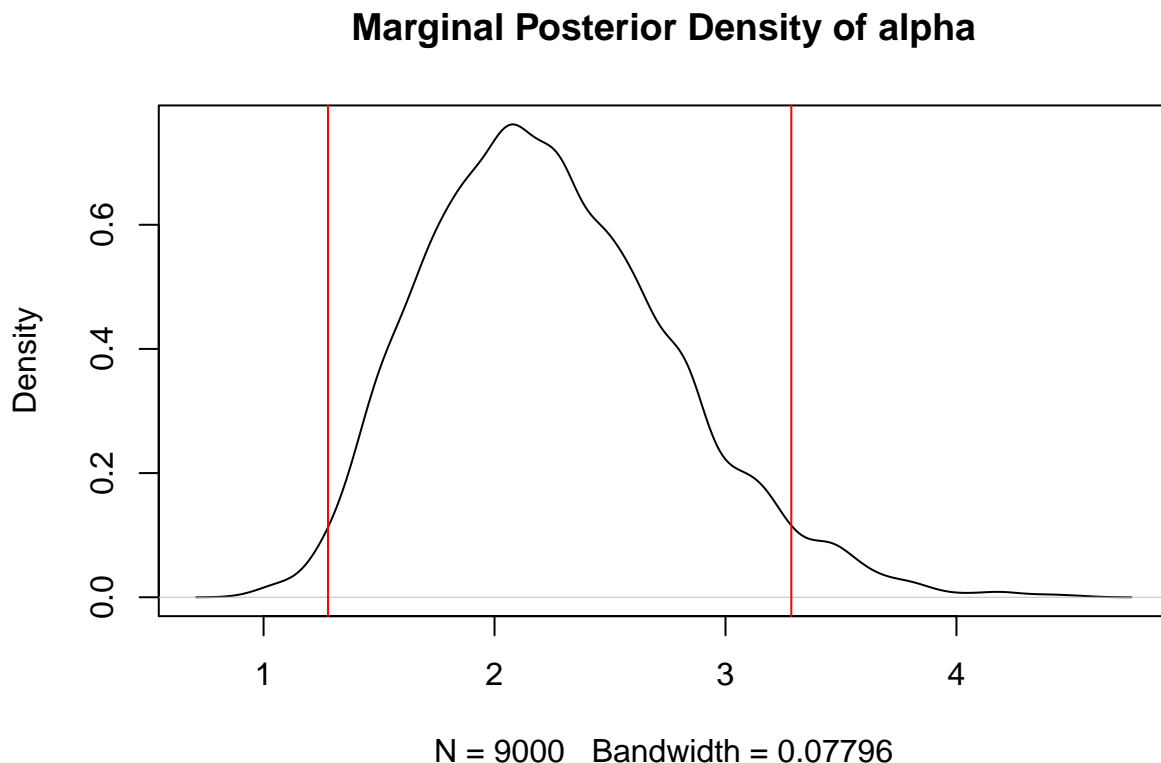
```
##  
##      Stationarity start    p-value  
##      test      iteration  
## var1 passed      1      0.585  
## var2 passed      1      0.551  
##  
##      Halfwidth Mean Halfwidth  
##      test  
## var1 passed      2.25 0.0701  
## var2 passed      2.16 0.0705
```

The chains pass the test of stationarity.

Marginal posterior density

Plot the MCMC estimate of the marginal posterior density for each parameter with an 95% HPD interval for each parameter

```
library(TeachingDemos)  
d1 <- density(mh.draws[,1])  
plot(d1, main="Marginal Posterior Density of alpha")  
abline(v=emp.hpd(mh.draws[,1], conf=0.95), col = "red")
```




```
emp.hpd(mh.draws[,1], conf=0.95) #95% HPD interval for alpha
```

```
## [1] 1.278886 3.284908
```

```
d2 <- density(mh.draws[,2])  
plot(d2, main="Marginal Posterior Density of beta")  
abline(v=emp.hpd(mh.draws[,2], conf=0.95), col = "red")
```

Marginal Posterior Density of beta



```
emp.hpd(mh.draws[,2], conf=0.95) #95% HPD interval for beta
```

```
## [1] 1.159514 3.306168
```